

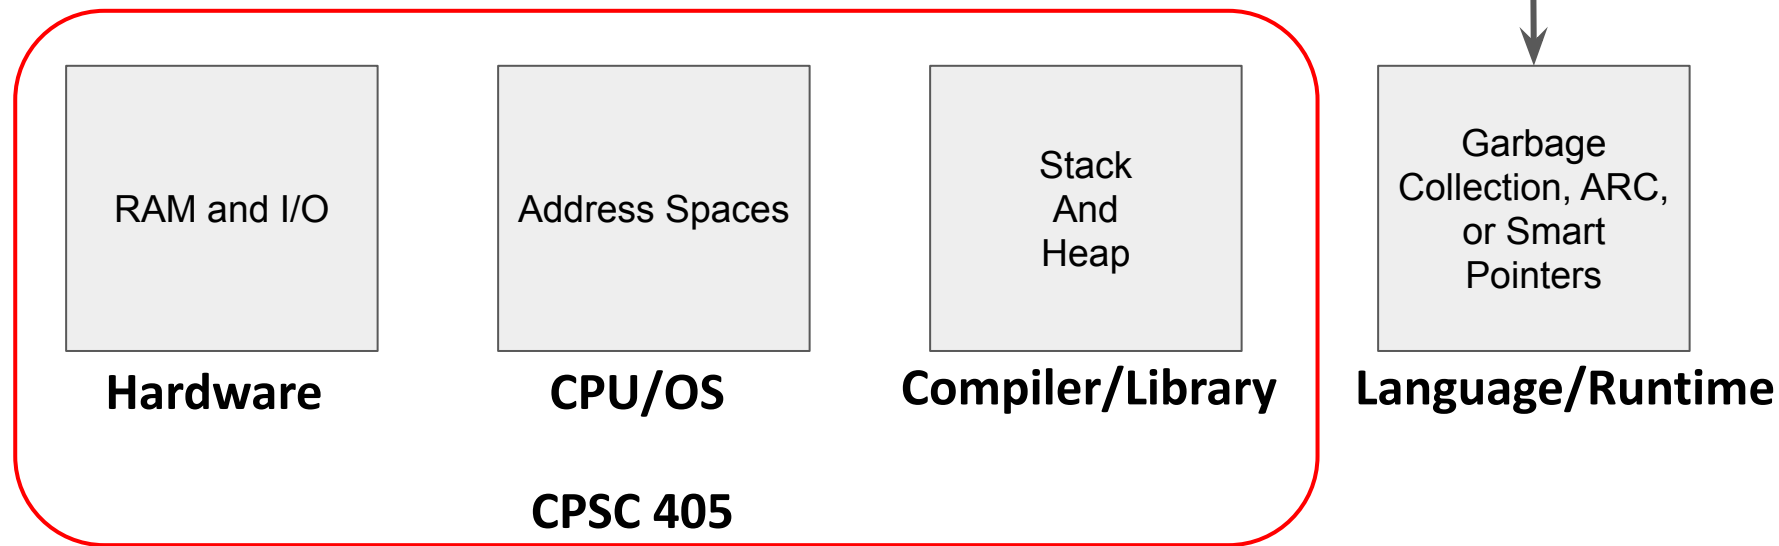
Programming Xv6 in C

Modified Charts from MIT's 6.181


Outline

- What is memory?
- C programming basics
- Logistics
 - Don't forget to post lecture questions before each lecture (starting this Wed.). • If you do so the night before, we'll try to cover it in lecture
 - The first lab is due this Thursday


Memory Abstractions



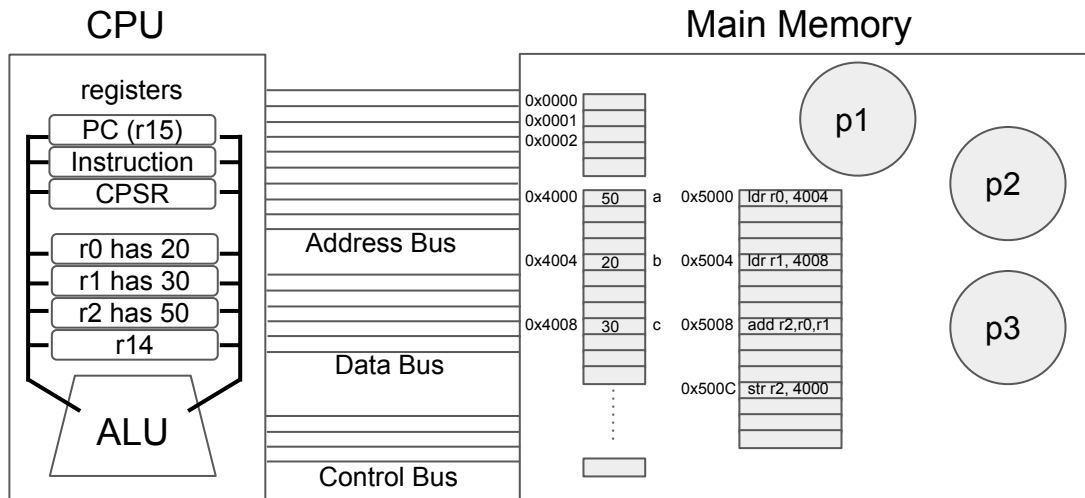
Hardware Layer



- A bus transfers data between components in the computer
- A cache remembers data previously fetched from the bus
- Speeds up the CPU by reducing the number of bus accesses
- Q: What is an IO Device?



Hardware Layer - Bus



OS shares
CPU among
processes


C Code

```
int a, b = 20, c = 30;  
a = b + c;
```

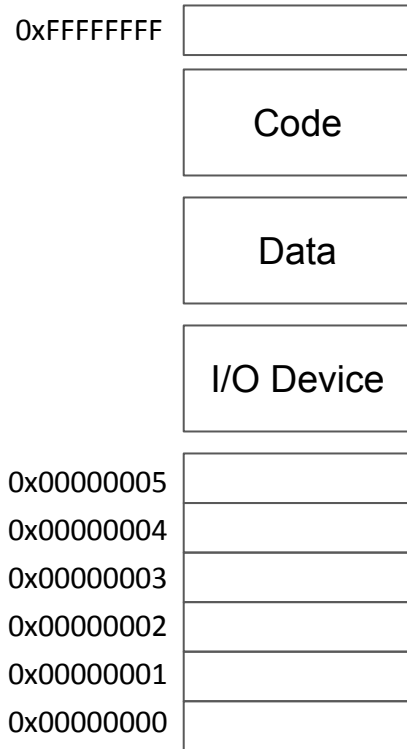
Assembly Code

```
ldr r0, #0x4004  
ldr r1, #0x4008  
add r2, r0, r1  
str r3, #0x4000
```

OS allocates
memory for
processes



Address Space - Array (or sequence) of bytes



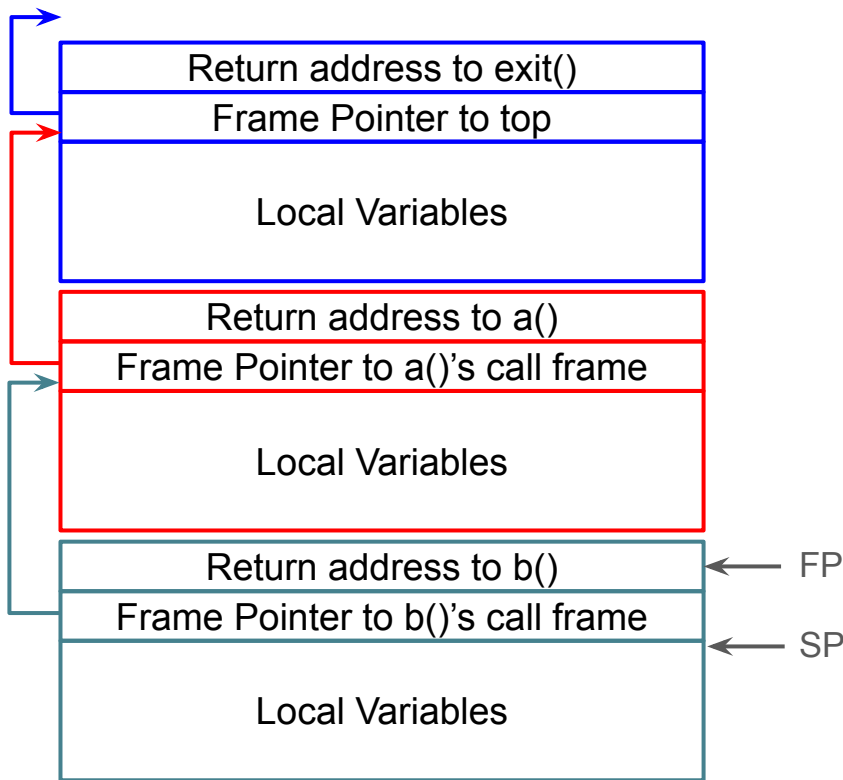
- Array index is address of a byte (8-bits)
- Loads and stores access address space
- Address space - virtual addresses, RAM - physical addresses
- Address space is usually much larger than RAM
- Addresses that can be accessed are referred to as “mapped”
- And holes that can’t be accessed are “unmapped”
- Address space has permissions: Read - load, Write - store, Execute - allow code to execute
- Address space combines RAM and I/O devices
- What happens if the CPU loads or stores to an unmapped region?
- Why have permissions?
- What happens if the CPU loads/stores an address without permission?

Stack Basics

High memory

Stack
Grows

Low memory



a(args...)
b(args...)
c(args...)

Heap Basics

- `void *malloc(size_t size)`
 - Allocates an object of size bytes
 - Returns 0 if out of memory! Otherwise, a pointer to the object.
- `void free(void *item);`
 - Frees an object
 - Can't be called more than once on same object
- Using heap API

```
struct foo *f = malloc(sizeof(*f));  
if (!f)  
    // handle out of memory error  
memset(f, 0, sizeof(*f)); // initialization  
// do something with f  
free(f);
```

Stack and Heap in an Address Space

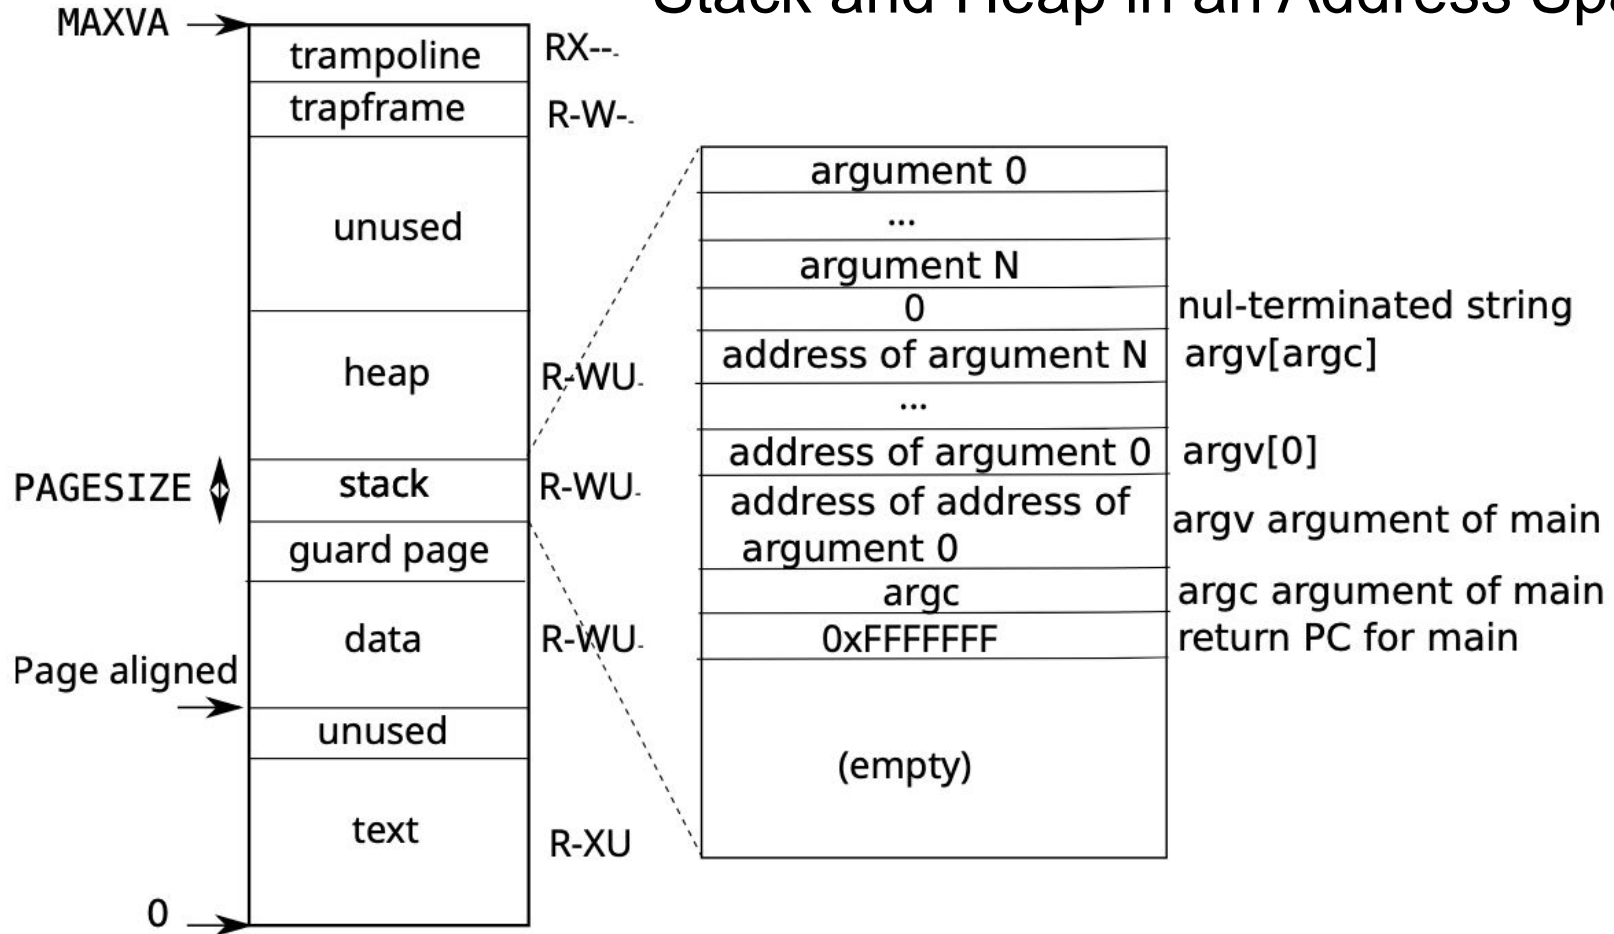


Figure 3-4 of Xv6 Book

Common Memory Problems

- Using memory after freeing it
- Freeing the same object more than once
- Forgetting to initialize memory (nothing is zeroed automatically)
- Writing beyond the end of an array (buffer overflow)
- Forgetting to free an object (memory leak)
- Casting an object to the wrong type
- Forgetting to check if an allocation failed
- Using pointers to locations on the stack (if they could return)

Why use C for an OS

- Good for low-level programming
 - Can manipulate address spaces directly without language abstractions
 - Easy to access hardware structures and RISC-V instructions
- Kernel is in complete control of memory allocation
 - In fact, you can build a memory allocator using C
 - No garbage collection
- Efficient and fast: compiled, no interpreter
- Why not?
 - Easy to write incorrect/insecure code!
 - Limited abstractions.

Primitive Types - RISC-V

- char: 1 byte
- short: 2 bytes
- int: 4 bytes
- long: 8 bytes
- long long: 8 bytes
- void *: 8 bytes (any pointer type is this size)
- Qualifiers: unsigned (nonnegative), const (can't be modified), static (only accessed within the file)
- sizeof(type) returns the size of a type
-

Using C typedefs

- xv6 uses typedefs to make the size of types more obvious
- `typedef unsigned char uint8; // uint8 is the same as unsigned char`
- `typedef unsigned short uint16;`
- `typedef unsigned int uint32;`
- `typedef unsigned long uint64;`

C structs

```
struct a {  
    int foo;  
};
```

```
struct b {  
    struct a bar;  
    long baz;  
};
```

Q: What will `printf("%ld", sizeof(struct b))` print?

Casting

- Converts one type to another
- Example:
 - `int foo = 10;`
 - `long bar = (long)foo;`

Pointer Arithmetic

```
void foo(void *ptr)
{
    void *pos = ptr + 10;      // doesn't compile!
    void *pos = (char *)ptr + 10; // works fine
    uint64 addr = (uint64)pos;  // can convert to int
    addr += 10;
    pos = (void *)addr;        // and back again
}
```

Bitwise Operators

`0b10001 & 0b10000 == 0b10000`

`0b10001 | 0b10000 == 0b10001`

`0b10001 ^ 0b10000 == 0b00001`

`~0b1000 == 0b0111`

Arrays

```
int foo[5];
```

```
int i;
```

```
for (i = 0; i < 5; i++) {
```

```
    foo[i] = i; }
```

```
// now foo contains 0, 1, 2, 3, 4
```

What does this print?

```
#include <stdio.h>
int main() {
    int x[5];
    printf("%p\n", x); // equivalent to &x[0]
    printf("%p\n", x+1); // equivalent to &x[0] + 1
    printf("%p\n", &x); // pointer to x
    printf("%p\n", &x+1); // eqv. to x + sizeof(x[5])
    return 0;
}
```

Source: <https://blogs.oracle.com/linux/post/the-ksplice-pointer-challenge>

Conclusion

- Many layers of abstraction in memory
- Writing an OS requires you to be aware of all of them
- C is a low-level language, so it's good at doing this
- But many pitfalls; large potential for bugs and security problems