

# Condition Variables

# Semaphores

# Concurrency Tools and Problems

- Concurrency Tools (Primitives)
  - Locks, Condition Variables, Semaphores
- Problems
  - Mutual exclusion
    - Threads A and B run separately in a critical region
    - Solved with locks
  - Ordering
    - Thread B runs after thread A (or A runs after B)
    - Solved with condition variables and semaphores

# Barrier Synchronization (ordering) Problem (part of Lab Multithreading)

- A group of threads where all threads must stop computing at a the barrier and wait until all of the group has arrived at the barrier








pthread Barrier API (from [https://en.wikipedia.org/wiki/Barrier\\_\(computer\\_science\)](https://en.wikipedia.org/wiki/Barrier_(computer_science)))

- `pthread_barrier_init()`  
Initialize the thread barrier with the number of threads needed to wait at the barrier in order to lift it
- `pthread_barrier_destroy()`  
Destroy the thread barrier to release back the resource
- `pthread_barrier_wait()`  
Calling this function will block the current thread until the number of threads specified by `pthread_barrier_init()` call `pthread_barrier_wait()` to lift the barrier.

# Producer-Consumer - Linux pipes - Ordering Problem

- A pipe may have many writers and readers
- Internally, there is a finite-sized buffer
- Writers (producers) add data to the buffer
  - Writers must wait if buffer is full
- Readers (consumers) remove data from the buffer
  - Readers must wait if buffer is empty
- Producer-Consumer or Bounded Buffer

# Bounded Buffer - Linux pipe

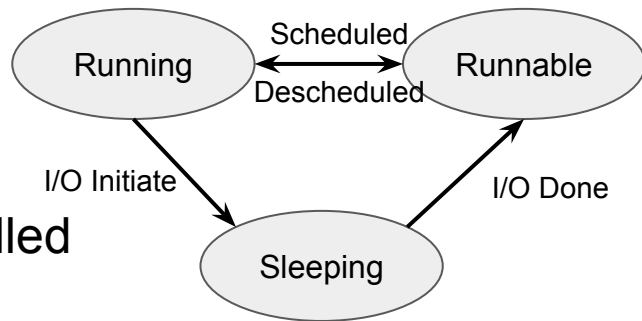


# Condition Variable - has queue of waiting threads

- Thread B waits for a signal on condition variable before running
  - `wait(CV, ...)` - Thread B sleeps (or blocks)
- Thread A signals condition variable when time for Thread B to run
  - `signal(CV, ...)` - Thread B moves from sleeping to runnable.

## Condition Variable API

- `wait(cond_t *cv, mutex_t *lock)`
  - Assumes the lock is held when `wait()` is called
  - Puts caller to sleep and releases the lock
  - When awoken, reacquires lock before returning
- `signal(cond_t *cv)`
  - Wake a single waiting thread (if  $\geq 1$  thread is waiting)
  - If there is no waiting thread, just return, doing nothing
    - The signal is not remembered, the signal is gone



# Producer-Consumer Helper Functions

```
int max; // variables and main
int loops;
int *buffer;
int use_ptr = 0;
int fill_ptr = 0;
int num_full = 0;
int main(int argc, char **argv) {
    max = atoi(argv[1]);
    loops = atoi(argv[2]);
    consumers = atoi(argv[3]);

    buffer = malloc(max*sizeof(int));
    ...
}
```


```
// buffer functions
void do_fill(int value) {
    buffer[fill_ptr] = value;
    fill_ptr = (fill_ptr + 1) % max;
    num_full++;
}

int do_get() {
    int tmp = buffer[use_ptr];
    use_ptr = (use_ptr + 1) % max;
    num_full--;
    return tmp;
}
```

# Producer-Consumer - Incorrect - One CV

```
void *producer(void *arg) {
    for (int i=0; i<loops; i++) {
        mutex_lock(&m);           //p1
        while(num_full == max)    //p2
            cond_wait(&cond, &m); //p3
        do_fill(i);              //p4
        cond_signal(&cond);       //p5
        mutex_unlock(&m);         //p6
    }
}
```

```
void *consumer(void *arg) {
    while(1) {
        mutex_lock(&m);           //c1
        while(num_full == 0      //c2
            cond_wait(&cond, &m); //c3
        int tmp = do_get();      //c4
        cond_signal(&cond);      //c5
        mutex_unlock(&m);        //c6
        printf("%d\n", tmp);     //c7
    }
}
```



does last signal wake producer or consumer2?



# Producer-Consumer - Correct - Two CVs

```
void *producer(void *arg) {  
    for (int i = 0; i < loops; i++) {  
        mutex_lock(&m);  
        while (numfull == max)  
            cond_wait(&empty, &m);  
        do_fill(i)  
        cond_signal(&fill);  
        mutex_unlock(&m);  
    }  
}
```

```
void *consumer(void *arg) {  
    while (1) {  
        mutex_lock(&m);  
        while (numfull == 0)  
            cond_wait(&fill, &m);  
        int tmp = do_get();  
        cond_signal(&empty);  
        mutex_unlock(&m);  
    }  
}
```

Correct!

- no concurrent access to shared state
- every time lock is acquired, assumptions are reevaluated
- a consumer will get to run after every do\_fill()
- a producer will get to run after every do\_get()

# Programming with Condition Variables

- Programming uses a condition variable, mutex, and state variable
  - Keep state (numfull in prior charts) in addition to CV's
- Must always do wait/signal with mutex lock held
- Whenever thread wakes from waiting, recheck state (This is the while loop)
  - Possible for another thread to grab lock in between signal and wakeup from wait

# Condition Variables and Semaphores

- Condition variables have a wait queue, but they do not have other state information
  - Programmer tracks state with variables
  - For example, we added the variable `numfull` for the producer/consumer solution
- Semaphores have have a wait queue and an integer state
  - State is maintained by the semaphore semantics. Calling semaphore API alters the underlying state

# Semaphore API

```
sem_t sem;  
sem_init(sem_t *s, int zero, int initval)  
sem_post(sem_t *s)  
sem_wait(sem_t *s)
```

- `sem_init` - initializes the integer state of the semaphore
- `sem_wait` - decrement the semaphore by 1, wait on the queue if the value is negative.
- `sem_post` - increment the semaphore by 1, if there are threads waiting on the queue, wake one.

# Create Mutex Lock with Semaphore

```
typedef struct __lock_t {  
    sem_t sem;  
} lock_t;
```

sem\_wait(): Decrement sem, If < 0, wait  
sem\_post(): Increment sem, then wake a waiter

```
void init(lock_t *lock) {  
    sem_init(&lock->sem, 1); // create sem with val == 1  
}  
void acquire(lock_t *lock) {  
    sem_wait(&lock->sem);  
}  
void release(lock_t *lock) {  
    sem_post(&lock->sem);  
}
```

# Create Semaphore with CV

```
typedef struct {    void sem_init(sem_t *s, int value) {
    int value;      s->value = value;
    cond_t cond;    cond_init(&s->cond);
    lock_t lock;    lock_init(&s->lock);
} sem_t;           }
```

sem\_wait(): Decrement sem, If < 0, wait  
sem\_post(): Increment sem, then wake a waiter

```
sem_wait(sem_t *s) {
    mutex_lock(&s->lock);
    while (s->value <= 0)
        cond_wait(&s->cond);
    s->value--;
    mutex_unlock(&s->lock);
}

sem_post(sem_t *s) {
    mutex_lock(&s->lock);
    s->value++;
    cond_signal(&s->cond);
    mutex_unlock(&s->lock);
}
```

See [zemaphore.c](#) in Lab Multithreading

# Producer/Consumer - Semaphores

- Circular Buffer - multiple producer threads, multiple consumer threads
- Shared buffer with N elements between producer and consumer
- Solution uses 2 semaphores
  - emptyBuffer: Initialize to N, producer can put N items in buffer
  - fullBuffer: Initialize to 0, consumer cannot consume until producer puts

Producer

```
int fill = 0;
put(int v) {
    buffer[fill] = v;
    fill = (fill+1) % N
```

```
}
```

```
i = 0;
```

```
while (1) {
```

```
    sem_wait(&emptyBuffer);
```

```
    put(value);
```

```
    sem_post(&fullBuffer);
```

```
}
```

Consumer

```
int use = 0;
```

```
int get() {
```

```
    int t = buffer[use];
```

```
    use = (use+1) % N;
```

```
    return t;
```

```
}
```

```
j = 0;
```

```
while (1) {
```

```
    sem_wait(&fullBuffer);
```

```
    int value = get();
```

```
    sem_post(&emptyBuffer);
```

```
}
```

must guard



critical region

# Producer/Consumer - Semaphores

- Circular Buffer - multiple producer threads, multiple consumer threads

Producer

```
int fill = 0;
put(int v) {
    buffer[fill] = v;
    fill = (fill+1)%N
}
```

```
sem_t mutex = 1;
i = 0;
```

```
while (1) {
```

```
    sem_take(&mutex);
```

```
    sem_take(&emptyBuffer);
```

```
    put(value);
```

```
    sem_post(&fullBuffer);
```

```
    sem_post(&mutex);
```

```
}
```

Consumer

```
int use = 0;
```

```
int get() {
```

```
    int t = buffer[use];
```

```
    use = (use+1)%N;
```

```
    return t;
```

```
}
```

```
j = 0;
```

```
while (1) {
```

```
    sem_take(&mutex);
```

```
    sem_take(&fullBuffer);
```

```
    int value = get();
```

```
    sem_post(&emptyBuffer);
```

```
    sem_post(&mutex);
```

```
}
```

Consumer waiting  
on fullBuffer

Producer waiting  
on mutex

Deadlock



# Producer/Consumer - Semaphores

- Circular Buffer - multiple producer threads, multiple consumer threads

Producer

```
int fill = 0;
put(int v) {
    buffer[fill] = v;
    fill = (fill+1)%N
}
sem_t mutex = 1;
i = 0;
while (1) {
    sem_take(&emptyBuffer);
    sem_take(&mutex);
    put(value);
    sem_post(&mutex);
    sem_post(&fullBuffer);
}
```

Consumer

```
int use = 0;
int get() {
    int t = buffer[use];
    use = (use+1)%N;
    return t;
}
j = 0;
while (1) {
    sem_take(&fullBuffer);
    sem_take(&mutex);
    int value = get();
    sem_post(&mutex);
    sem_post(&emptyBuffer);
}
```

Deadlock Fixed

# Semaphores

- Semaphores are equivalent to locks and condition variables
  - Can be used for both mutual exclusion and ordering
- Semaphores contain state - an integer
  - How they are initialized depends on how they will be used
  - Init to 1: Mutex
  - Init to N: Number of available resources - producer/consumer solution

## Semaphore API

- `sem_wait` - decrement the semaphore by 1, wait if the value  $< 0$
- `sem_wait` - waits until value  $> 0$ , then decrement (atomic)
- `sem_post` - increment the semaphore by 1, if there are threads waiting, wake one

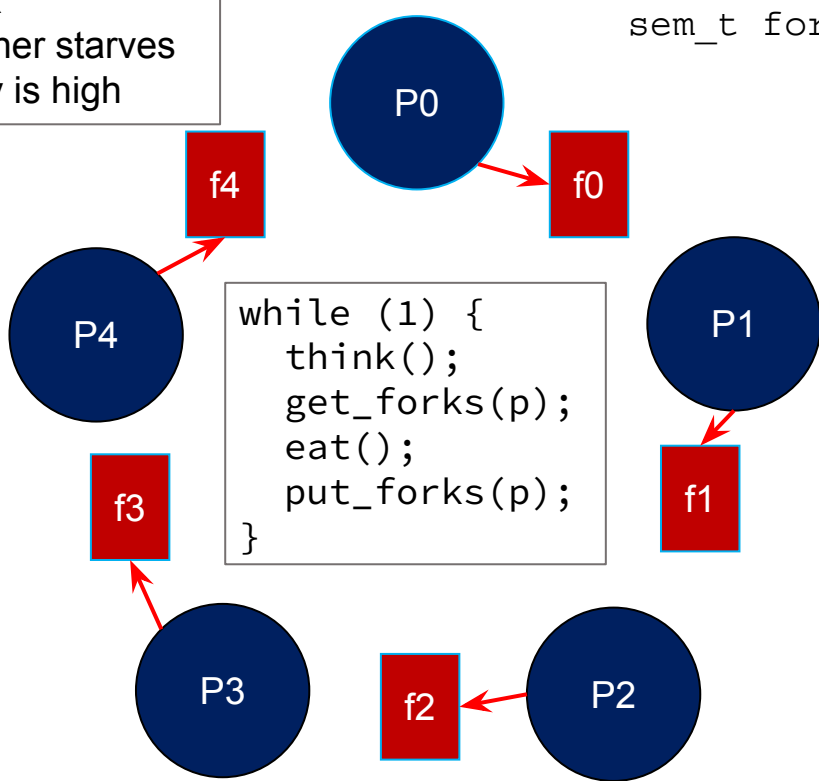
# Producer-Consumer Code

- Lab Multithreading (git checkout thread), folder notxv6/pc has code
  - pc\_cv.c - condition variable solution to producer-consumer
  - pc\_sem.c - semaphore solution to producer-consumer
- Run both programs with various values for
  - <buffer size> <loops> <consumers>
  - Write observations in your notebook
- Update the programs to allow for multiple producers, run your updated program
  - Write observations in your notebook
- Questions
  - How is the buffer defined?
  - What is the argument passed to producer and consumer threads?
  - What is the technique used to terminate the consumer threads?
  - Which of the two solutions do you like better? Why?

# Dining Philosophers - (See OSTEP Ch 31, pp 13-18)

Write `get_forks()` and `put_forks()` such that

- no deadlock
- no philosopher starves
- concurrency is high



## Deadlock Solution

Philosophers all grab forks to the left.  
Everyone is waiting on a fork.

```
void get_forks(int p) {  
    Sem_wait(&forks[left(p)]);  
    Sem_wait(&forks[right(p)]);  
}  
20  
  
void put_forks(int p) {  
    Sem_post(&forks[left(p)]);  
    Sem_post(&forks[right(p)]);  
}
```

# Dining Philosophers - (See OSTEP Ch 31, pp 13-18)

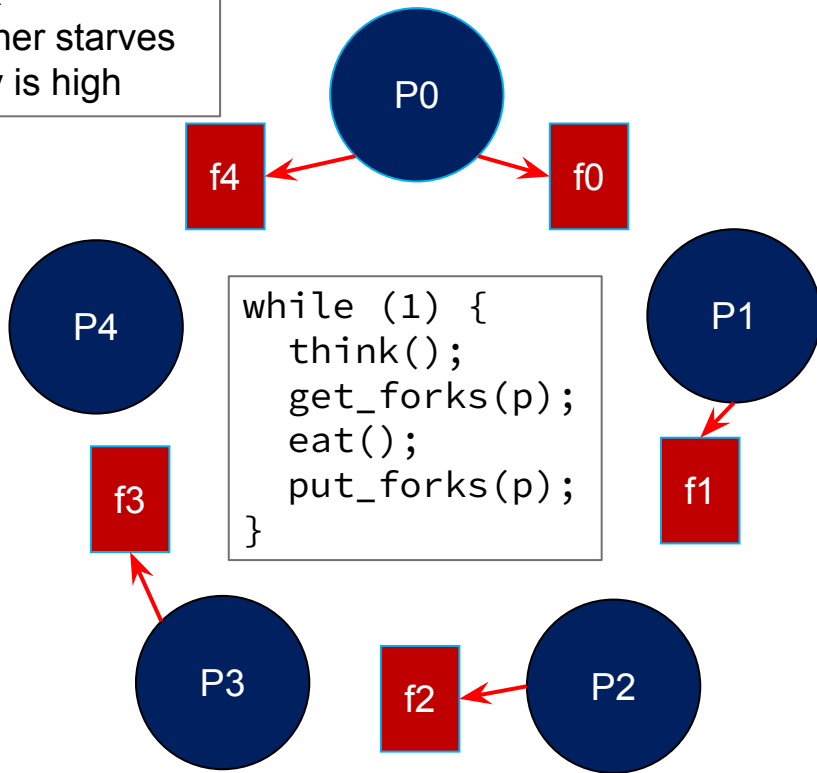
Write `get_forks()` and `put_forks()` such that

- no deadlock
- no philosopher starves
- concurrency is high

```
sem_t forks[5] = {1};
```

Correct Solution

Four Philosophers all grab lowest number forks first.  
Fifth philosopher grabs highest number fork first.



```
void get_forks(int p) {  
    if (p == 0) {  
        Sem_wait(&forks[right(p)]);  
        Sem_wait(&forks[left(p)]);  
    } else {  
        Sem_wait(&forks[left(p)]);  
        Sem_wait(&forks[right(p)]);  
    }  
}  
  
void put_forks(int p) {  
    Sem_post(&forks[left(p)]);  
    Sem_post(&forks[right(p)]);  
}
```