

File Systems

Part 1

Charts: Augmented from MIT's Adam Belay

<https://pdos.csail.mit.edu/6.828/2023/lec/l-fs1.pdf>

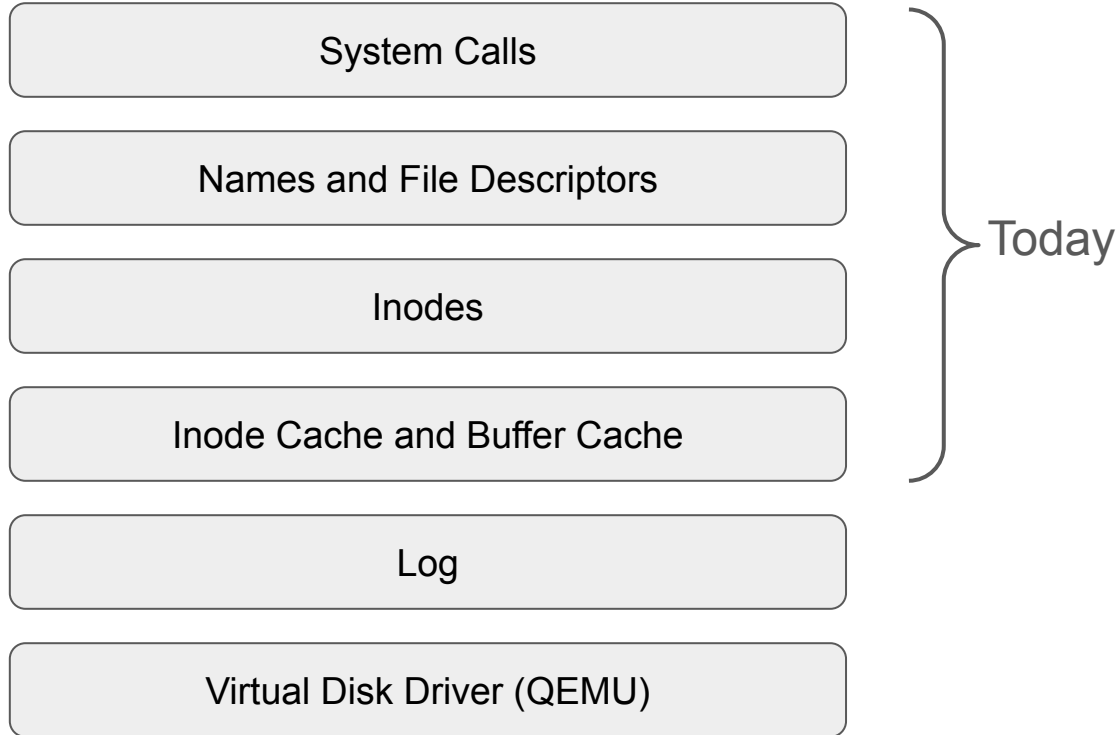
File Systems

- Provide persistent storage across restarts and crashes
- Provide naming and organization
- Provide sharing of data among users and processes

File System Attributes that are interesting to study/solve

- Crash recovery
- Performance + concurrency
- Sharing + security
- Powerful abstractions (e.g., proc, afs, 9P, pipes, etc.)

File System Software Layers



High-level design choices in system calls

- Objects: Use files (not virtual disks or databases)
- Content: Use byte arrays (not structured)
- Naming: Human-readable (not ID numbers)
- Organization: Name hierarchy
- Synchronization: None (no locking, no versions)
- `link()` and `unlink()` can change names concurrently w/ `open()`

System Call Layer

```
fd = open("x/y", flags); // creates a
fd
write(fd, "abc", 3);      // wr 3 bytes
link("x/y", "x/z");      // create link
unlink("x/y");            // removes x/y
write(fd, "def", 3);      // wr 3 more
bytes
close(fd);                // close the
fd
```

File x/z contains abcdef
See file user/user.h.

```
int pipe(int*);
int write(int, const void*, int);
int read(int, void*, int);
int close(int);
int kill(int);
int exec(char*, char**);
int open(const char*, int);
int mknod(const char*, short, short);
int unlink(const char*);
int fstat(int fd, struct stat*);
int link(const char*, const char*);
int mkdir(const char*);
int chdir(const char*);
int dup(int);
```

System Calls

Names and File
Descriptors

Inodes

Inode Cache and
Buffer Cache

Log

Virtual Disk Driver
(QEMU)

Linux File System Call API

https://www.gnu.org/software/libc/manual/html_node/File-System-Interface.html

```
int fd = open(char *path, int flag, mode_t mode)
int close(int fd)
ssize_t read(int fd, void *buf, size_t nbyte)
ssize_t write(int fd, void *buf, size_t nbyte)
int fsync(int fd)
off_t lseek(int fd, off_t offset, int whence)
int rename(const char *old_filename, const char *new_filename)
int stat(const char *path, struct stat *statbuf)
int fstat(int fd, struct stat *statbuf)
int dup(int old) - copies old to first available fd
int dup2(int old, int new) - copies old to new, replacing new
int link(const char *oname, const char *nname)
int symlink(const char *oname, const char *nname)
int unlink(const char *pathname)
int mkdir(const char *path, mode_t mode)
int chdir(const char *path)
int fchdir(int filedes)
char *getcwd(char *buffer, size_t size)
DIR *opendir(char *dirname)
struct dirent *readdir(DIR *dirp)
```

System Calls

Names and File
Descriptors


Inodes

Inode Cache and
Buffer Cache

Log

Virtual Disk Driver
(QEMU)

Name Layer



- Path names are organized as a tree
- No cycles, but multiple names can refer to the same file (i.e., via `link()`)
- Processes share the namespace
- But each process has a current working directory (CWD)
- Absolute path: `/x/y`
- Relative path: `x/y`

System Calls

Names and File
Descriptors

Inodes

Inode Cache and
Buffer Cache

Log

Virtual Disk Driver
(QEMU)

File Description Layer

- Each process has its own FD number namespace
- Each FD identifies a file created by `open()`
- By convention `STDIN`, `STDOUT`, `STDERR` are file descriptors 0, 1, and 2
- Lowest available FD number is allocated during `open()`
- FD returned during `open()` is usable by the process even if the file is unlinked (i.e., deleted)
- A file is an object that you can read and write to like a stream

System Calls

Names and File
Descriptors


Inodes

Inode Cache and
Buffer Cache

Log


Virtual Disk Driver
(QEMU) 8

Interacting with a file



- FDs access file as an array of bytes, very similar to an address space
- Each FD has a “**cursor**” to the file
- An empty file **cursor** starts at byte 0

Interacting with a file




- FDs access file as an array of bytes, very similar to an address space
- Each FD has a “**cursor**” to the file
- `read()` advances the **cursor**

Interacting with a file



- FDs access file as an array of bytes, very similar to an address space
- Each FD has a “**cursor**” to the file
- `read()` advances the **cursor**
- `write()` advances the **cursor**

Pipe Files are a bounded buffer



- read blocks when the buffer is empty
- write blocks when the buffer is full
- Multiple processes can read/write from/to a pipe

Inode Layer

- **Inode:** Records the details of a file
 - Tracks the size of the file and where on the disk the data is stored
 - Has a link count (and open FD count) to figure out when to free
 - Deallocation deferred until link + open count is zero
- **I-number:** Refers to an inode, similar to an FD
 - Uniquely identifies a position on disk where Inode is located

System Calls

Names and File
Descriptors

Inodes

Inode Cache and
Buffer Cache

Log

Virtual Disk Driver₁₃
(QEMU)

Where is Data Stored

- On a persistent storage medium
- Persistent - data doesn't go away under loss of power
- Common storage mediums
 - HDDs: High capacity, slow, inexpensive
 - SSDs: Lower capacity, faster, more expensive
- Disks accessed in fixed-sized units (like pages)
 - Called sectors, historically 512 bytes
 - Newer drives use 4K (4096) byte sectors

Performance Characteristics

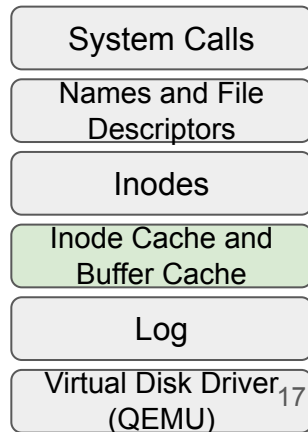
- Applies to both HDDs and SSDs, but more so the HDDs
- Sequential access much faster than random
- Big reads/writes much faster than small ones
- Both facts influence FS design

Disk Blocks


- Disk blocks are allocated to files
- Typically, multiple sectors are combined to form blocks
- e.g., a 4KB block is 8 512-byte sectors
- Combining sectors helps reduce book-keeping and seek overhead
- Xv6 uses two 512-byte sector blocks
- Every block has a block number
- Block number is like an address that identifies the location on the disk

Inode Cache and Buffer Cache Layer

- Disk accesses are slow and random
- Store copies of inodes and blocks in RAM
- Works well because the same data is often accessed many times
- e.g., the same inodes and blocks are accessed each time a file is read
- No need to access the disk if a copy is available!



Xv6 Disk Layout



- Xv6 provides mkfs program
- Generates this layout for a new (empty) FS
- The layout is static for the lifetime of the FS
- mkfs macros
 - `#define NINODES 200`
 - `#define BSIZE 1024`
 - `#define FSSIZE 2000` ← blocks on disk
 - `#define IPB (BSIZE / sizeof(struct dinode))`
 - `int nbitmap = FSSIZE/(BSIZE*8) + 1;`
 - `int ninodeblocks = NINODES / IPB + 1;`
 - `int nlog = LOGSIZE;`
- File System Metadata
 - Everything other than file content
 - Super block, inodes, bitmap, directory content

Inodes

Disk Inode

```
struct dinode { // disk inode
    short type;    // File type: dir, file
    short major;   //
    short minor;   //
    short nlink;   // # links to inode
    uint size;     // file sz (bytes)
    uint addrs[NDIRECT+1]; // Data blocks
}
```

NDIRECT is 12

sizeof(struct dinode) is **64 bytes**

$2 + 2 + 2 + 2 + 4 + (13 * 4)$

16 dinodes per disk block

Memory Inode

```
struct inode { // mem copy of inode
    uint dev;    // device number
    uint inum;   // inode number
    uint ref;    // refs to inode
    struct sleeplock lock // prot all below
    int valid;   // inode read from disk
    short type;  // File type: dir, file
    short major; //
    short minor; //
    short nlink; // # links to inode
    uint size;   // file sz (bytes)
    uint addrs[NDIRECT+1]; // Data blocks
}
```

System Calls

Names and File
Descriptors

Inodes

Inode Cache and
Buffer Cache

Log


Virtual Disk Driver
(QEMU)¹⁹

- What are some file attributes missing from these inodes?

int fd = open("file.txt", ,);

```
struct proc {
    struct spinlock lock;
    // p->lock held for these
    enum procstate state;          // Process state
    void *chan;                    // sleep channel
    int killed;                    // !=0, killed
    int xstate;                    // exit status returned to parent
    int pid;                       // Process ID
    // wait_lock held for this
    struct proc *parent;           // Parent proc
    // p->lock not needed for these private
    uint64 kstack;                 // kernel stack for proc
    uint64 sz;                     // Size proc mem (bytes)
    pagetable_t pagetable;         // user pagetable
    struct trapframe *trapframe;   // data for trampoline
    struct context *context;        // switch()
    struct file *ofiles[NFILE];    // Open files
    struct inode *cwd;             // Current dir
    char name[16];                 // Process name
};
```

```
struct dirent {
    ushort inum;
    char name[DIRSIZ];
};
```



```
struct file {
    enum {FD_NONE, FD_PIPE,
          FD_INODE, FD_DEVICE } type;
    int ref;                       // ref count
    char readable;
    char writable;
    struct pipe *pipe;             // FD_PIPE
    struct inode *ip;           // FD_INODE
    uint off;                      // rd/wr spot
    short major;                   // FD_DEVICE
};
```

```
struct inode { // mem copy of inode
    uint dev;                     // device number
    uint inum;                    // inode number
    uint ref;                     // refs to inode
    struct sleeplock lock;        // prot all below
    int valid;                    // inode read from disk
    short type;                   // File type: dir, file
    short major;                 //
    short minor;                 //
    short nlink;                 // # links to inode
    uint size;                    // file sz (bytes)
    uint addrs[NDIRECT+1];       // Data blocks
};
```

On Disk Inode Layout

- type: Free, file, directory, or device
- nlink: number of links
- size: the size of the file in bytes
- addrs: addresses of data blocks (array)
- Example: Find file's byte at 4000
 - $4000 / \text{BSIZE} (=1024) = 3$;
 - `addrs[3]` contains the block number (659) with data byte at 4000


```
struct dinode { // disk inode
    short type; // File type: dir, file
    short major; //
    short minor; //
    short nlink; // # links to inode
    uint size; // file sz (bytes)
    uint addrs[NDIRECT+1]; // Data blocks
}
```

NDIRECT is 12

`sizeof(struct dinode)` is **64 bytes**

$2 + 2 + 2 + 2 + 4 + (13 * 4)$

16 dinodes per disk block



Inode Fixed Size - File Sizes

- How can we fit large files into addrs?
- Use indirect block: a full block of more addrs
- Xv6 has one indirect block
- You can let addrs be all indirect blocks

```
struct dinode { // disk inode
    short type; // File type: dir, file
    short major; //
    short minor; //
    short nlink; // # links to inode
    uint size; // file sz (bytes)
    uint addrs[NDIRECT+1]; // Data blocks
}
```


NDIRECT is 12

sizeof(struct dinode) is **64 bytes**



$2 + 2 + 2 + 2 + 4 + (13 * 4)$

16 dinodes per disk block

Indirect
Block: 73



... 256 additional blocks



Converting an Inum to an Inode

- Two-hundred Dinodes on disk blocks 32 through 44
- i-number is used to select a Dinode
- Each Dinode is 64 bytes long
- If you read the Dinode blocks into memory
- Byte Offset of Inode(i-number) = $64 * i\text{-number}$
- Directories contain data struct for files
- Directory data struc has i-numbers
- Use i-num to get Dinode
- Use Dinode to get file
- We use both Dinodes and Inodes

Block 44	Offset 12288
Inode 192	


Inode 191	Offset 2048
Block 34	
Inode 32	

Inode 31	Offset 1024
Block 33	
Inode 16	

Inode 15	Offset 0
Block 32	
Inode 0	

Directories

- A directory is a file.
- Directory has a name and blocks with data
- Users can not directly write contents
- Content is an array of dirents. Each direct:
 - i-number (of the file in the directory)
 - 14-byte file name (Xv6 filenames are old-school)
 - dirent is free (unused) if inum == 0
 - sizeof(struct dirent) is 16 bytes
 - A 1024 byte block holds 64 struct dirent



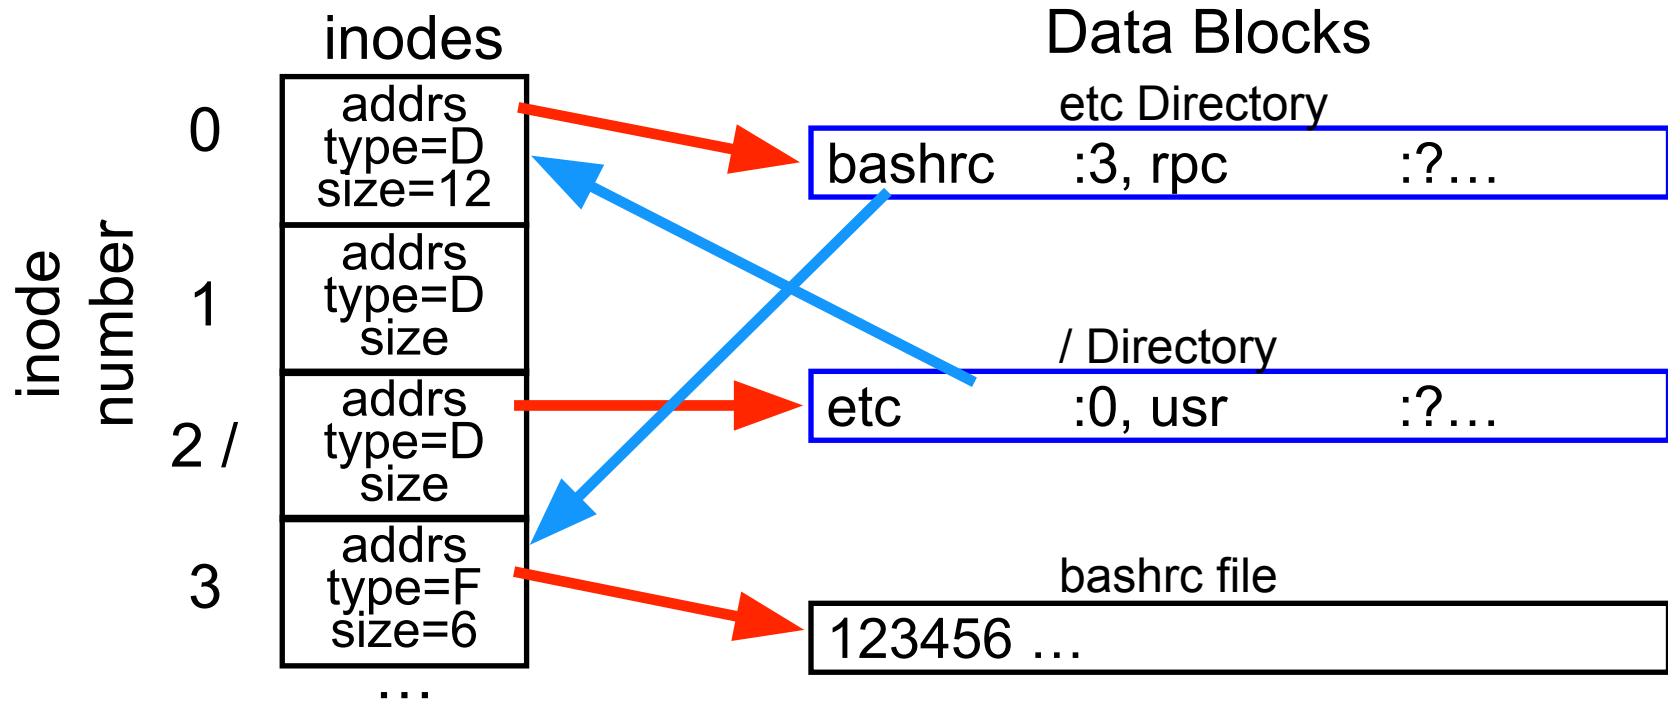
On Disk Structure is a Tree

- Layer 1: Directory
- Layer 2: Inodes
- Layer 3: Disk Blocks
- Layer 4: Disk Sectors

Pools of information to Allocate

- Inodes
- Blocks

read /etc/bashrc



Reads for traversing /etc/bashrc

read root dir (inode and data)
read etc dir (inode and data)
read bashrc file (inode and data)

```
struct dirent {  
    ushort inum;  
    char name[DIRSIZ];  
};
```