

File Systems

Part 2

Charts: Augmented from MIT's Adam Belay

<https://pdos.csail.mit.edu/6.828/2023/lec/l-fs2.pdf>

Crash Recovery

- **Issue:** Crashes leave disk in inconsistent state
- **Solution:** Logging

Crash Scenario

- Imagine you are using the filesystem
- Power is suddenly lost
- The system reboots
- Is the filesystem still usable?
- Is your data still there?

Crash Scenario is a Difficult Problem

- Filesystems perform multi-step operations
 - reserve an inode, then reserve a bit in the bitmap, then update a directory, then fill in an inode, then write data, etc.
 - A crash could leave invariants violated
- After rebooting:
 - Bad outcome: Crash again due to corrupt FS
 - Worse outcome: Silently read/write invalid data

create /foo/bar

data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data
		read			read	
			read			read
	read write					write
				read write		
			write			

```

struct dinode { // disk inode
    short type; // File type: dir, file
    short major; //
    short minor; //
    short nlink; // # links to inode
    uint size; // file sz (bytes)
    uint addrs[NDIRECT+1]; // Data blocks
}
    
```

```

struct dirent {
    ushort inum;
    char name[DIRSIZ];
};
    
```

open /foo/bar

data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data
		read			read		
			read				
				read			
						read	

```
struct dinode { // disk inode
    short type; // File type: dir, file
    short major; //
    short minor; //
    short nlink; // # links to inode
    uint size; // file sz (bytes)
    uint addrs[NDIRECT+1]; // Data blocks
}
```

```
struct dirent {
    ushort inum;
    char name[DIRSIZ];
};
```

write to /foo/bar (assume file exists and has been opened)

data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data
read write				read			
				write			write

```
struct dinode { // disk inode
    short type; // File type: dir, file
    short major; //
    short minor; //
    short nlink; // # links to inode
    uint size; // file sz (bytes)
    uint addrs[NDIRECT+1]; // Data blocks
}
```

```
struct dirent {
    ushort inum;
    char name[DIRSIZ];
};
```

read /foo/bar – assume opened

data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data
				read			
				write			
							read

```
struct dinode { // disk inode
    short type; // File type: dir, file
    short major; //
    short minor; //
    short nlink; // # links to inode
    uint size; // file sz (bytes)
    uint addrs[NDIRECT+1]; // Data blocks
}
```

```
struct dirent {
    ushort inum;
    char name[DIRSIZ];
};
```

close /foo/bar

data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data

```
struct dinode { // disk inode
    short type; // File type: dir, file
    short major; //
    short minor; //
    short nlink; // # links to inode
    uint size; // file sz (bytes)
    uint addrs[NDIRECT+1]; // Data blocks
}
```

nothing to do on disk!

```
struct dirent {
    ushort inum;
    char name[DIRSIZ];
};
```


Xv6 Various disk writes for file creation/population

- `$ echo hi > x`
- `bwrite`: block 33 by `ialloc` // allocate inode (block 33)
- `bwrite`: block 33 by `iupdate` // update inode (e.g., set `nlink`)
- `bwrite`: block 46 by `writei` // write directory entry
- `bwrite`: block 32 by `iupdate` // update directory inode with new `len`

```
struct dirent {  
    ushort inum;  
    char  
    name[14];  
};
```

```
struct dinode {  
    short type;  
    short major;  
    short minor;  
    short nlink;  
    uint size;  
    uint addrs[12+1];  
}
```

create /x				
data bitmap	inode alloc	root inode	x inode	root data
read		read		read
		write write		write write

46: Data Blocks

45: Free Blk BM

32: Inodes

3: Log Blocks

2: Log Head

1: Super Block

Various disk writes for file creation/population

- `$ echo hi > x`
- `bwrite: block 33 by ialloc // allocate inode (block 33)`
- `bwrite: block 33 by iupdate // update inode (e.g., set nlink)`
- `bwrite: block 46 by writei // write directory entry`
- `bwrite: block 32 by iupdate // update directory inode with new len`

Crash →

What Happens

- Not much bad happens
- Inode allocated and wasted, never usable in future

What about this order?

- `$ echo hi > x`
- `bwrite: block 46 by writei // write directory entry`
- `bwrite: block 32 by iupdate // update directory inode with new len`
- `bwrite: block 33 by ialloc // allocate inode (block 33)`
- `bwrite: block 33 by iupdate // update inode (e.g., set nlink)`

Crash →

What Happens

- Disaster!
- Inode could be reallocated again
- Directory points to uninitialized inode

What Order Could Really Happen

- Kernel (and maybe the disk too) reorders writes to minimize seeks
- In general, any order is possible

Writing Files Has Multiple Disk Writes

1. inode `addrs[]` and `len`
2. indirect block
3. block content
4. block free bitmap

Crash Scenario 1: inode refers to free block -- disaster!

Crash Scenario 2: block not free but not used -- not so bad

Unlink has Multiple Disk Writes

1. block free bitmaps
2. free inode
3. erase `dirent`

File System Recovery Goals

After reboot, run recovery code

1. Internal FS invariants must hold
 - e.g., no block is both free and used by an inode
2. All but the last few operations stored on disk
 - Data I wrote yesterday should be there!
 - But perhaps data at the time of crash will be lost
3. No reordering of data writes
 - `echo 99 > result ; echo done > status`

Correctness and performance

- Often at odds with one another!
- Disk writes are very slow
- Safety: Write data right away
- Speed: Wait and batch together writes

Crash recovery

- Arises in all storage systems, e.g., fs, databases
- Many clever solutions
- Performance/correctness tradeoffs

Logging (or Journaling)

- Goal: Atomic system calls w.r.t. Crashes
- Goal: Fast recovery (no hour-long fsck)
- xv6: minimal design for safety
- ext3: adds more speed

Logging basics

- Atomicity: All of system call's writes applied or none are
- Each atomic op is called a transaction
- Three phase operation:
 - 1. Log phase: Record all the writes the system call will perform on disk
 - 2. Commit phase: Record done on disk
 - 3. Install phase: Do the actual disk writes

Crash recovery w/ logging

- Crash recovery of complex mutable data structures is hard
- But logging makes it easy, can retrofit on top of existing FS designs
- If "done" found in log, replay all writes
- If "done" not found, ignore entries in log
- Called write-ahead logging

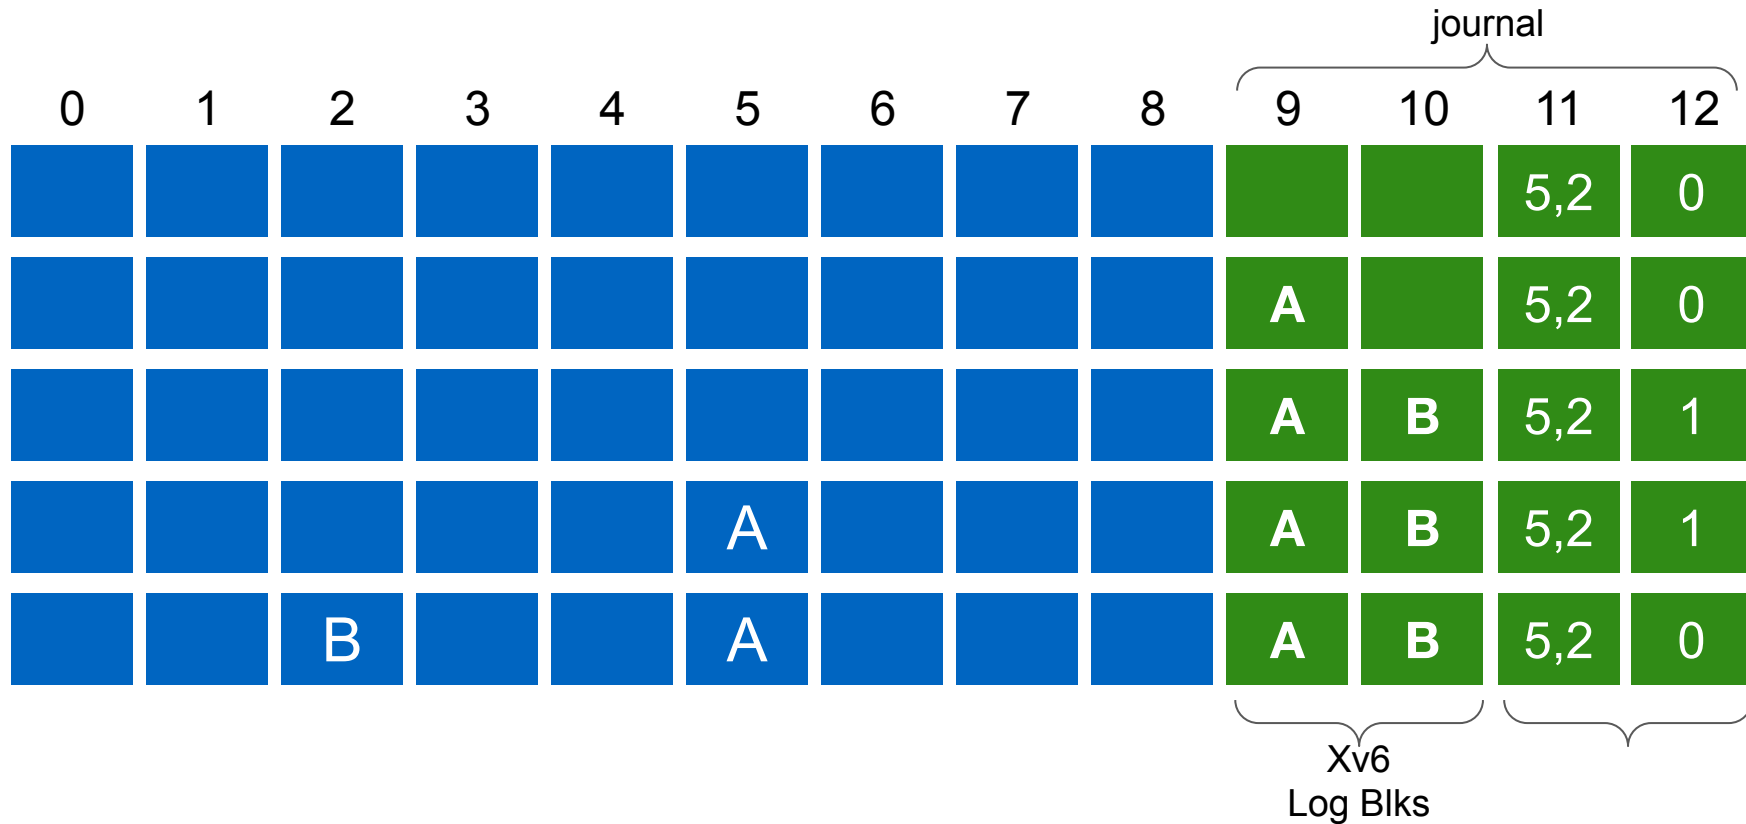
Rules

- install *none* of a transaction's writes to disk
- until *all* writes are in the log on disk, and the logged writes are marked committed

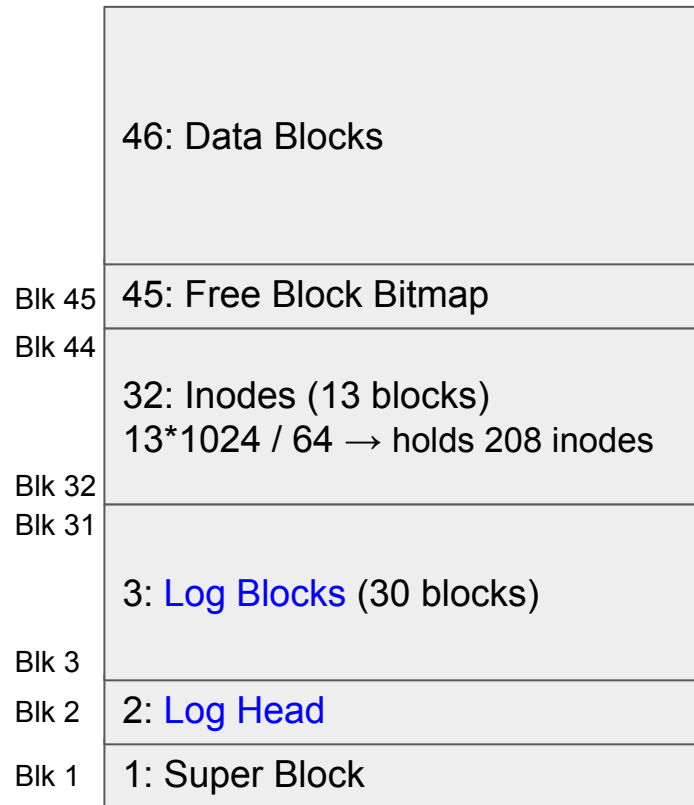

Once we've installed a transaction on disk...

- We have to do all its writes
- This ensures the transaction is atomic
- Log allows us to detect if all steps in the transaction are there
- If not, we can safely abort the transaction

transaction: write A to block 5; write B to block 2



Xv6 Logging



Xv6 Logging Steps

- On write:
 - Add blockno to in-memory array
 - Keep the data itself in the buffer cache (pinned)
- On commit:
 - Write buffered log to disk
 - Wait for disk to complete writing (synchronous)
 - Write the log header sector to disk
- After commit:
 - Install (write) the blocks in the log to their location in FS
 - Unpin the blocks in the buffer cache
 - Write zero to the log header sector on disk
- NOTE - X6 assumes
 - Either an entire sector is written or it is not (no partial writes)
 - No decay of sectors (no read errors)
 - No read of the wrong sector (seek errors)

Log Header


- An “n” value on disk indicates the commit point
- Nonzero: Indicates a valid transaction is committed on disk
- Zero: Not committed, may not be complete
- Records block #s that were updated
- And the number of blocks in log

\$ echo “hi” > x

- Create inode
- Write ‘hi’ to file x
- Write ‘\n’ to file x


Create File X

- bwrite 3 // inode, 33
- bwrite 4 // directory content, 46
- bwrite 5 // directory inode, 32
- bwrite 2 // commit (block #s and n)
- bwrite 33 // install inode for x
- bwrite 46 // install directory content
- bwrite 32 // install dir inode
- bwrite 2 // mark log "empty"




Write “hi” to x

- bwrite 3 // bitmap update (45)
- bwrite 4 // actual data (746)
- bwrite 5 // inode update (33)
- bwrite 2 // commit (block #s and n)
- bwrite 45 // bitmap
- bwrite 746 // a (note: bzero was absorbed)
- bwrite 33 // inode (file size)
- bwrite 2 // mark log "empty"



Write “\n” to x

- bwrite 3 // actual data (746)
- bwrite 4 // inode update (33)
- bwrite 2 // commit (block #s and n)
- bwrite 746 // \n
- bwrite 33 // inode (file size)
- bwrite 2 // mark log "empty"



Logging Challenges

- Challenge: Prevent write-back
 - Buffer cache holds in-memory copies of disk blocks
 - Can't let the buffer cache write back until logged
 - Tricky because cache could run out of memory
- xv6's solution:
 - Ensure buffer cache is big enough
 - Pin dirty blocks in buffer cache (can't cycle out)
 - After commit, unpin blocks
- Challenge: Data must fit in the Log
- xv6 solution:
 - Compute upper bound on number of blocks each system call could write
 - Set the log size to be greater than this upper bound
 - Break up some system calls into several transactions
 - E.g., really large write()'s
 - Large writes not atomic, but crash will leave correct prefix

Summary

- Logging makes file system transactions atomic
 - Either they complete fully or not at all
- Write-ahead logging is the key solution in xv6
 - Log written in batches, good for performance
 - But now each disk write happens twice!
 - And have to wait (synchronous) for disk writes
 - Trouble with operations that don't fit in log
- Overall, performance is quite a bit worse
 - Next lecture: How can we make this fast?