

Threads Locks

Synchronization Primitives

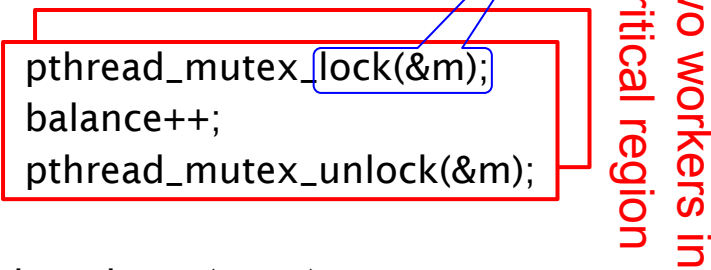
- Locks or Mutex or Mutex Lock
 - Used to guard a critical region
 - Guarantees mutual exclusion to the critical region
- Condition Variables and Semaphores
 - Used to order the execution of threads. Guarantee that thread Y waits until thread X finishes
 - Semaphores can be used for mutual exclusion and for ordering

pThread Code

```
pthread_mutex_t m;  
volatile int balance = 0;  
int loops = 2000;  
void *worker(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        pthread_mutex_lock(&m);  
        balance++;  
        pthread_mutex_unlock(&m);  
    }  
    pthread_exit(NULL);  
}
```

Only one thread
at a time

Two workers in
critical region



```
int main(int argc, char *argv[]) {  
    pthread_t p1, p2;  
    printf("Initial value : %d\n", balance);  
    pthread_mutex_init(&m, NULL);  
    pthread_create(&p1, NULL, worker, NULL);  
    pthread_create(&p2, NULL, worker, NULL);  
    pthread_join(p1, NULL);  
    pthread_join(p2, NULL);  
    printf("Final value : %d\n", balance);  
    return 0;  
}
```

Critical Region

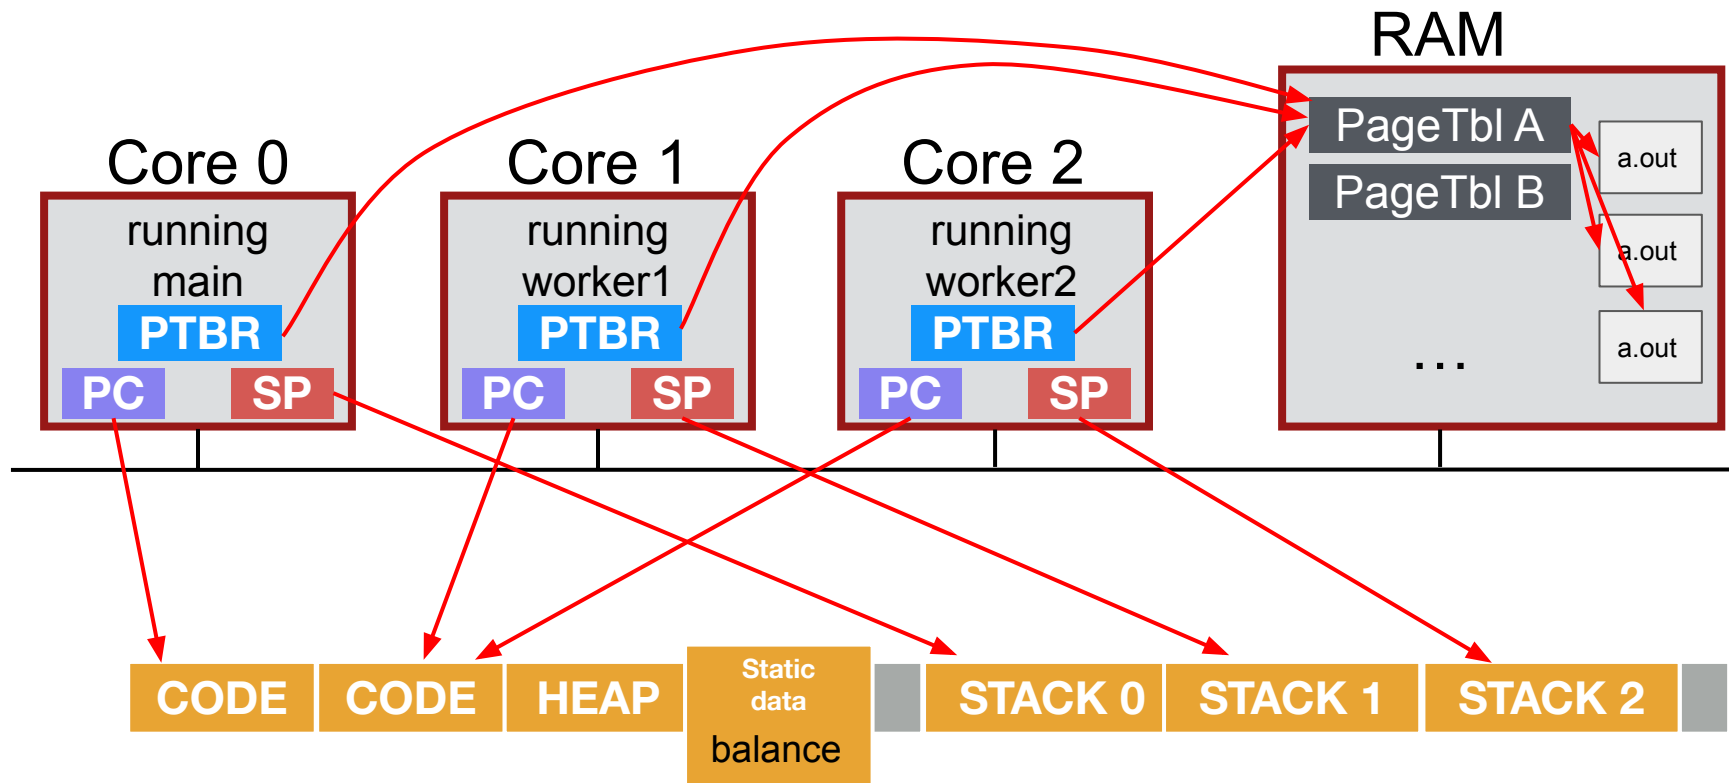
- When two threads update the same static variables. The update of one thread must be complete before the second thread. C code `balance++` → critical region
- The assembly instructions corresponding to the C code must be completed without getting interrupted

`balance++`

```
lla a5,balance
lw  a5,0(a5)
addiw a5,a5,1
sxt.w a4,a5
lla a5,balance
sw  a4,0(a5)
```

→ critical region

- Each thread must have mutual exclusive access to critical regions



```
Pthread_mutex_lock(&m);  
    balance++;  
Pthread_mutex_unlock(&m);
```

Locking Algorithm - Incorrect

```
int lock = 0; // 0 means free
```

```
void acquire(int *lock) {  
    while (*lock); ← spin till lock == 0  
    *lock = 1;  
}
```

```
void release(int *lock) {  
    *lock = 0;  
}
```

```
release:  
    addi sp,sp,-32  
    sd s0,24(sp)  
    addi s0,sp,32  
    sd a0,-24(s0)  
    ld a5,-24(s0)  
    sw zero,0(a5)  
    ld s0,24(sp)  
    addi sp,sp,32  
    jr ra
```

```
acquire:  
    addi sp,sp,-32  
    sd s0,24(sp)  
    addi s0,sp,32  
    sd a0,-24(s0)  
    .L2:  
    ld a5,-24(s0)  
    lw a5,0(a5)  
    bne a5,zero,.L2  
    ld a5,-24(s0)  
    li a4,1  
    sw a4,0(a5)  
    ld s0,24(sp)  
    addi sp,sp,32  
    jr ra
```

Race Condition with Load and Store

*lock == 0

Thread 1

while(*lock)

interrupted here

*lock = 1

acquire:

addi sp,sp,-32

sd s0,24(sp)

addi s0,sp,32

sd a0,-24(s0)

.L2:

ld a5,-24(s0)

lw a5,0(a5)

bne a5,zero,.L2

ld a5,-24(s0)

li a4,1

sw a4,0(a5)

ld s0,24(sp)

addi sp,sp,32

jr ra

Thread 2

while(*lock)

*lock = 1

Both threads grab lock!
Testing lock and setting
lock are not atomic

Atomic Assembly Instructions

```
.var mutex
.var balance
.main
.top
```

```
for (i = loop; i > 0; i++) {
    lock(mutex);
    balance = balance + 1;
    unlock(mutex)
}
```

```
.acquire
```

```
mov  mutex, %ax      # loop until mutex is 0
test $0, %ax         # we can be interrupted in the loop
jne  .acquire
mov   $1, %ax         # maybe another thread got mutex as 0
xchg %ax, mutex      # atomic swap of 1 (in ax) and mutex
test $0, %ax         # if we get 0 back: lock is free!
jne  .acquire        # if not, try again
```

Spin Loop

Thread Snuck In

```
# critical section - must have lock to enter
mov  balance, %ax     # get the value at the address
add  $1, %ax          # increment it
mov  %ax, balance     # store it back
```

Critical Region

```
# release lock
mov  $0, mutex
# see if we're still looping
sub  $1, %bx          # reg bx has i
test $0, %bx
jgt  .top
halt
```


Spin Lock Implementation with xchg

```
typedef struct lock {  
    int locked;  
} lock;  
  
void init(lock *lk) {  
    lk->locked = 0;  
}  
  
void acquire(lock *lk) {  
    while(xchg(&lk->locked, 1));  
}  
  
void release(lock *lk) {  
    lk->locked = 0;  
}
```

```
int xchg(int *adr, int new) {  
    int old = *adr;  
    *adr = new;  
    return old;  
}
```

```
xchg(volatile unsigned int *addr, unsigned int newval) {  
    uint result;  
    asm volatile("lock; xchgl %0, %1" : "+m"  
                (*addr), "=a" (result) : "1" (newval) : "cc");  
    return result;  
}
```

Lock Algorithm - Wrong

```
struct lock { int locked; }  
acquire(l) {  
    while(1){  
        if(l->locked == 0){ // A  
            l->locked = 1;    // B  
            return;  
        }  
    }  
}
```

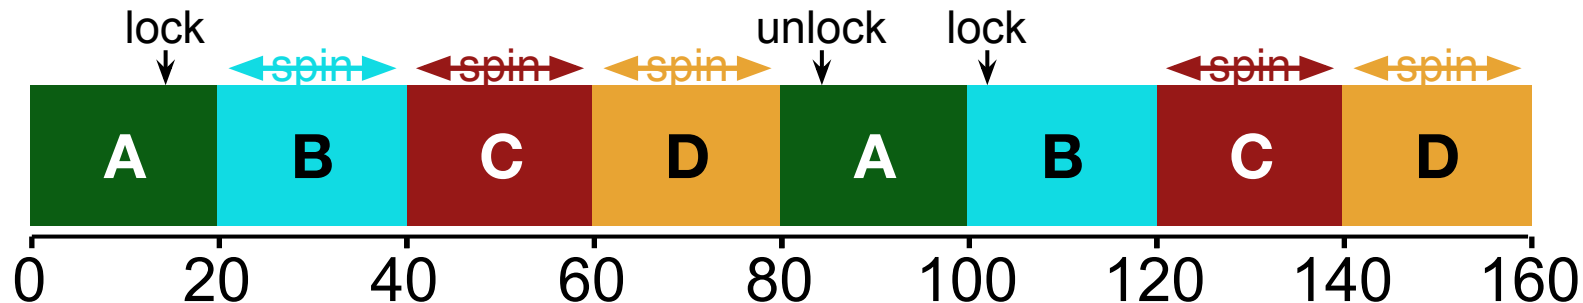
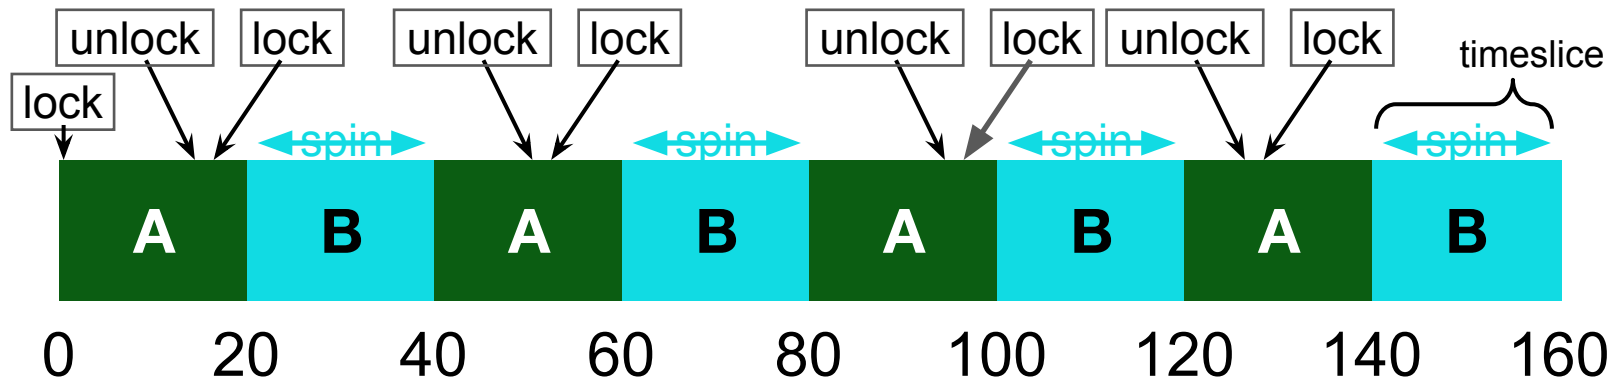
oops: race between lines A and B
how can we do A and B atomically?

Xv6 Spin Lock Implementation

```
acquire(lk){  
    while(__sync_lock_test_and_set(&lk->locked, 1) != 0)  
}
```

- if lk->locked was already 1, sync_lock_test_and_set sets to 1 (again), returns 1, and the loop continues to spin
- if lk->locked was 0, at most one lock_test_and_set will see the 0; it will set it to 1 and return 0; other test_and_set will return 1
- this is a "spin lock", since waiting cores "spin" in acquire loop

Spinlocks are not fair - threads wait equally




Scheduler is not aware of locks/unlocks

Spin Lock Implementation with `yield`

```
typedef struct lock {
    int locked;
} lock;
void init(lock *lk) {
    lk->locked = 0;
}
void acquire(lock *lk) {
    while(xchg(&lk->locked, 1))
        yield();
}
void release(lock *lk) {
    lk->locked = 0;
}
```

```
int xchg(int *adr, int new) {
    int old = *adr;
    *adr = new;
    return old;
}
```

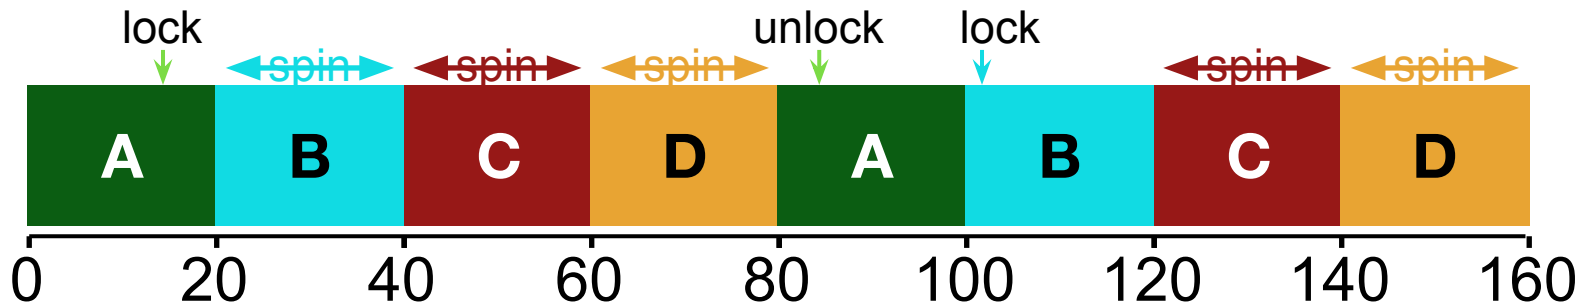
Yield Instead of Spin




Spinlock Performance

- Fast when...
 - many CPUs
 - locks held a short time
 - advantage: avoid context switch
- Slow when...
 - one CPU
 - locks held a long time
 - Without yield: $O(\text{num_threads} * \text{time_slice})$
 - With yield: $O(\text{num_threads} * \text{context_switch_time})$
 - disadvantage: spinning is wasteful
- **Next:** Block thread on queue instead of spinning


no yield:



yield:



block:



Waiting on Mutex - Spin vs Blocking

- Uniprocessor - thread waiting on mutex is scheduled
 - Blocking is better - does not waste CPU cycles
 - Associate queue of waiters with each lock
- Multiprocessor - thread waiting on mutex is scheduled
 - Spin or block depends on how long, t , before lock is released (proc spins) and the time, C , of a context switch
 - if $t < C$ (i.e., lock released quickly) - spinning is better
 - if $t > C$ (i.e., lock released slowly) - blocking is better
 - Quick and slow are relative to the length of time, t , a lock is held and context-switch time, C
 - You must understand your system design and implementation to know length a lock is held is

Spinning vs Blocking Analysis

- If we know how lock held time, we can determine optimal behavior
- CPU time wasted when spin-waiting is t , time spent spinning
- CPU time wasted when blocking is C , time of a context switch
- **Optimal**: spin time t is less than context time C ($t < C$), it is best to spin
- **Optimal**: spin time t is greater than context time C ($t > C$), it is best to block
- **Optimal**: achieved by knowing time lock will be held, which we don't know
- **Algorithm**: Spin-wait for C then block. This is 2 times the optimal
- **Case 1**: $t < C$: optimal spin-waits for t ; Algorithm spin-waits t too
- **Case 2**: $t > C$: optimal blocks immediately (cost of C); we pay spin C then block (cost of 2 C); $2C / C$, yields a algorithm that is 2 times the optimal

Concurrency Tools and Problems

- Concurrency Tools (Primitives)
 - Locks, Condition Variables, Semaphores
- Problems
 - Mutual exclusion
 - Threads A and B run separately in a critical region
 - Solved with locks
 - Ordering
 - Thread B runs after thread A (or A runs after B)
 - Solved with condition variables and semaphores